

COSC122 (2020) Lab 6.2

Sorting - $n \log n$ Algorithms

Goals

This lab will give you some experience with $O(n \log n)$ sorting algorithms; in this lab you will:

- implement a median-of-three partitioning algorithm;
- implement a version of quicksort to sort certain ranges of a list; and
- use it to quickly find the median value of a list.
- review merge sort and its *merge* operation.
- use a variant of the *merge* principle to find the common items on two lists.

You should be familiar with the material in Sections 5.35 to 5.36 of the textbook¹ before attempting this lab.

Quick sort

The `quicksort` module contains a number of quicksort methods, stubs, and helpers:

- `quicksort`: Starts the recursive quicksort process for sorting a list. It creates a copy of the list, so that the original isn't modified by `quicksort_helper` (which sorts in-place). The `quicksort` function returns a new sorted list - and doesn't affect the original.
- `quicksort_helper`: Recursively partitions and sorts a list between the `left` and `right` indices.
- `pivot_index`: Returns the index of the pivot point to use (at the moment, it simply returns the left-most value).
- `partition`: Partitions the section of a list between the `left` and `right` indices. Uses `left-pivot` or `mo3-pivot` - see below for details.
- `pivot_index_mo3`: Returns the index of the median value of three items from the list. The three items are the left, right and middle.
- `quicksort_range`: Starts the recursive quicksort process for a specific range within a list.
- `quicksort_range_helper`: Recursively partitions and sorts a specific range within a list (you will implement this later in the lab).

The basic quicksort implementation is complete, and you can test it in the shell:

```
>>> from quicksort import quicksort
>>> alist = [4,6,3,9,2,3,1]
>>> quicksort(alist)
[1, 2, 3, 3, 4, 6, 9]
```

¹Online text: see the Mergesort(5.11) and Quicksort(5.12) sections.

Median-of-Three Quick sort

Choosing a good pivot point is crucial to obtaining good quicksort performance. Ideally, you want to partition the list into two sub-lists that are equal in length—which means choosing a pivot that will be sorted to the middle of the list. However, you also want your pivot selection algorithm to be fast ($\mathcal{O}(1)$), so you can't spend a lot of time looking around the list for a good value.

In the examples of quicksort you've seen so far, the pivot has always been either the first or last item in the list. If the list of items is randomly arranged (the best case), then this is a perfectly adequate pivot point; however, if our list already happens to be sorted (the worst case), then this pivot would create highly unbalanced sub-lists—degrading quicksort to $\mathcal{O}(n^2)$.

A common approach to improving quicksort's partitioning algorithm is known as the *median-of-three* partition: we examine the first, middle, and last elements of the list and pick the median value of these three as our pivot.

As we are finding the median of three items we can simply sort them (using insertion sort makes sense here) and then the middle item is the median. For example, 3 is the median of the following lists [1, 3, 6], [3, 1, 6], [1, 6, 3], [3, 1, 3] and 2 is the median of the following lists [2, 2, 3], [2, 2, 2], [3, 1, 2], [2, 2, 1].

The provided quicksort uses the left item as a pivot by default (eg, using `s = quicksort(alist)` implicitly uses 'left-pivot'). If quicksort is called using 'mo3-pivot' then the partition function will call the `pivot_index_mo3` function to find the index of the pivot point to use—`pivot_i`.

Your job is to implement the `pivot_index_mo3` function and test the performance of quicksort when using it to find the pivot index (eg, `s = quicksort(alist, 'mo3-pivot')`).

DocTests are provided so that you can test your median of three function. You can comment out functions you don't want checked to avoid getting masses of irrelevant errors (eg, comment out `quicksort_range` at this stage).

Remember that `pivot_index_mo3` is returning the index of the median value, not the median value. For example, `pivot_index_mo3([5, 4, 3, 2, 12, 14, 10], 0, 4)` should return 0. This call is asking for the index of the item that is the median of the item at index 0, the item at index 4, and the item at index 2 (ie, the middle item). The three values to find the median for are [5, 3, 12] and therefore the median value is 5. The value 5 is at index 0 in the original list and therefore the function should return 0.

In Wing ctrl1- . can be used to toggle commenting of block of text.²

Add your own doctest(s) to the quicksort function to ensure that quicksort still works when called with 'mo3-pivot'.

> **Complete the Median of Three questions in Lab Quiz 6.2.**

Timing the Difference

The new pivot selection method doesn't change the big-O performance complexity of the quicksort implementation (for lists of random items), but it still gives us a speed-up when sorting *real-world* data.

The `quicksort_trials.py` file provides a graph³ of the time taken to sort ascending value lists of various lengths, using the default left-pivot quicksort.

You should now add in trials for random lists of values and graph them on the same graph - you can add any number of extra series to the plot function (eg, `pyplot.plot(x, y, 'bo', x1, y1, 'go')` plots two series; one with blue o's and one with green o's - 'bx' would do blue crosses, etc)⁴. The pyplot window may come up in the background and you should also remember to close the pyplot window if you want your program to finish.

²At least in Windoze and Linux.

³Lab computers have matplotlib installed. Check out the install guide(s) on the Quiz Server if you want to install them on your own computer—assuming you haven't already installed them for an earlier lab. Ask a tutor if you are having trouble.

⁴Marker styles use two chars with first being colour and second being type - feel free to experiment. Or, check the interwebs for information on pyplot marker styles

You can generate sorted lists with `range` as per the trial that is provided. Random lists can be generated using the `random.shuffle(mylist)` function:

```
test_list = list(range(800)) # test_list is a sorted list
random.shuffle(test_list) # mixes items in to a random order
```

Generate and print some sorted and random lists if you are unsure what is happening...

Note: Running quicksort with lists of greater than about 900 items will hit Python's internal recursion limit when lists are close to the worst case. Try it and see...

How does the time for sorted lists compare with the time for random lists?

Now, generate graphs for quicksort with the 'mo3-pivot' setting and compare the speeds. You should be able to graph all four runs on one graph, ie, sorted/random vs left-pivot/mo3-pivot. These graphs should give a feel for how sort time is related to the number of items in the list and the pivot method.

> **Complete the Quicksort Speed questions in Lab Quiz 6.2.**

Once you have finished playing with the graphs you should fill out Table 1 with the average time taken to sort 760 items in the four situations given. To calculate the time you should set the number of trials to 100, that is run each sort one hundred times and take the average time. When calculating ratios for the lab quiz you should try to do it in Python so that you don't get any rounding issues (or you can use average times that are written to 5 or 6 decimal places in your ratios).

	Sort time for random data n=760	Sort time for ascending data n=760
Left-Pivot Quicksort		
Median-of-three Quicksort		

Table 1: Your results from various quicksort runs with n=760 using 100 trials.

> **Complete the comparing pivot choices questions in Lab Quiz 6.2.**

Sorting only a limited range of the full list

Sometimes, you won't want to sort the *entire* list of items, but only a small range of items. For example, say we only want to see the first 10 search results (eg, the first page) from a Google search that returns 1,000,000 results — that is, the items in indices 0 to 9. We want the items in indices 0 to 9 to be sorted in the order they would be in if we sorted the entire list. This means that the items for that range can come from any part of the list, and they will be placed in such a way as if the whole list was sorted. For example, if you were to sort the range 8 to 10 of the list [10, 80, 90, 50, 60, 40, 20, 5, 100, 70, 30] you should receive something similar to: [5, 10, 30, 50, 60, 40, 20, 70, 80, 90, 100]. You will notice that the last three items [80, 90, 100] are sorted correctly and in the right position (ie, they are in the three top slots), but the rest of the list isn't quite sorted. Similarly if we started with the same original list but only wanted the smallest value we would use a start and end of 0 and the result would be that the first item in the list would be in order but the rest of the list wouldn't necessarily be—you would get something like [5, 10, 90, 50, 60, 40, 20, 80, 100, 70, 30]

To limit the range of indices that we want sorted, we'll pass in two extra parameters to our `quicksort` method: `start` and `end`, which specify the indices of the range to sort. The algorithm is almost the same as the `quicksort` method—it chooses a pivot, and partitions the list around it; but it will only do this if it needs to, ie, if the `start—end` range overlaps with the `left—right` range — for example, as in Figure 1.



Figure 1: A case when `quicksort_range_helper` does need to do something.

It may be easier to think of this in terms of when you should *NOT* recursively sort, shown in Figure 2.




Figure 2: Two cases when `quicksort_range_helper` doesn't need to do anything.

Implement the `quicksort_range_helper` method to sort a specified range of values, using a copy of the `quicksort_helper` code as a starting point and adjusting the base case condition so that irrelevant ranges are not sorted any further. `quicksort_range_helper` should ensure that items between the `start` and `end` index (inclusive) are fully sorted, but doesn't have to ensure that items outside the range are fully sorted.

`quicksort_range_helper` can use the same `partition` function.

`quick_sort_range` can still be run with `'left-pivot'` or `'mo3-pivot'`.

The following definition of median will help you answer the quiz questions below.

In statistics, the median M of a quantitative data set of size n is the middle number when the values are arranged in ascending (or descending) order. If n is odd, M is the middle number; if n is even, M is the mean of the middle two numbers.

> **Complete Quicksort Range questions in Lab Quiz 6.2.**

Mergesort and its merge

The main part of mergesort is as follows:

```
if the list has more than one item:
    split into left and right halves
    mergesort the left half
    mergesort the right half
    merge the two sorted halves
    return the merged list
else:
    return the list as a 0 or 1 item list is already sorted
```

The merge operation is where most of the heavy lifting is done. A simple version of the merge algorithm is given below. See the text book for a variation that merges the result back into the original list, effectively doing an in place sort⁵.

```
result = []
i=0
j=0
while i < len(left) and j < len(right):
    if left[i] <= right[j]:
        result.append(left[i])
        i=i+1
    else:
        result.append(right[j])
        j=j+1
# add any left-overs
while i < len(left):
    result.append(left[i])
    i=i+1
while j < len(right):
    result.append(right[j])
    j=j+1
```

> *Complete the Mergesort questions in Lab Quiz 6.2.*

Finding common items

Suppose you have been given the job of finding students who worked for the volunteer army after both the September and February earthquakes, so you can send them a token of thanks from the Cadbury chocolate company. You have a list of student ids from each army and you want a method of quickly listing ids that are in both lists.

The `common.py` module provides the stub for a `common_items(list_x, list_y)` function. Your job is to implement a twist on the *merge* algorithm (ie, with two indices that keep track of the current smallest item in each list) to generate the list of items that are common to both lists. You should include each common item only once in the output list. Doctests are provided to help you check your implementation.

NOTE: To start with we suggest you write code that can deal with the lists of unique items and then think about how to deal with lists that don't contain unique items.

To help you check that your routine works we provide some ordered files, in a folder called `data`, and a function for generating a list of integers from a given file, `read_data(filename)`. To access files in the `data` folder you should prefix them with `./data/`, for example `filename = './data/ordered_0.txt'`. There are two sets of files, one that contains lists of unique values and the other that doesn't — can you tell which is which?

> *Complete the Common items questions in Lab Quiz 6.2.*

⁵Textbook mergesort section.

Extra Exercises

Implement a `quick_median` function to use the `quicksort_range` method to find the median of an unsorted list. You should call `quicksort_range` with the appropriate start and end values (remember that the range we are interested in is different for odd-sized lists than that of even-sized lists).

Test the method to see how much faster it is for finding the median value, than sorting the *entire* list and extracting the median.